

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
```

```
In [2]: df = pd.read_csv('D:/Data Science/Scaler/PortfolioProjects/insurance.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Age	Diabetes	BloodPressureProblems	AnyTransplants	AnyChronicDiseases	Height	Weight	PremiumPrice
0	45	0	0	0	0	155	57	1327
1	60	1	0	0	0	180	73	2269
2	36	1	1	0	0	158	59	1699
3	52	1	1	0	1	183	93	3517
4	38	0	0	0	1	166	88	3151

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 986 entries, 0 to 985
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    986 non-null    int64
1   Diabetes                              986 non-null    int64
2   BloodPressureProblems                 986 non-null    int64
3   AnyTransplants                        986 non-null    int64
4   AnyChronicDiseases                    986 non-null    int64
5   Height                                986 non-null    int64
6   Weight                                986 non-null    int64
7   KnownAllergies                        986 non-null    int64
8   HistoryOfCancerInFamily               986 non-null    int64
9   NumberOfMajorSurgeries                 986 non-null    int64
10  PremiumPrice                           986 non-null    int64
dtypes: int64(11)
memory usage: 84.9 KB
```

Feature Engineering

```
In [5]: df['BMI'] = round(df['Weight'] / ((df['Height'] / 100) ** 2), 2)
```

```
In [6]: def categorize_bmi(BMI):
        if BMI < 18.5:
            return 'Underweight'
        elif BMI < 25:
            return 'Normal weight'
        elif BMI < 30:
            return 'Overweight'
        else:
            return 'Obese'

df['BMI_Category'] = df['BMI'].apply(categorize_bmi)
```

```
In [7]: df['Age_BMI'] = df['Age'] * df['BMI']
df['Chronic_Diabetes'] = df['AnyChronicDiseases'] * df['Diabetes']
df['BMI_Surgeries'] = df['BMI'] * df['NumberOfMajorSurgeries']
```

```
In [8]: df.head()
```


```
Out[8]:
```

	Age	Diabetes	BloodPressureProblems	AnyTransplants	AnyChronicDiseases	Height	Weight	...
0	45	0	0	0	0	155	57	...
1	60	1	0	0	0	180	73	...
2	36	1	1	0	0	158	59	...
3	52	1	1	0	1	183	93	...
4	38	0	0	0	1	166	88	...


BMI Analysis

```
In [9]: sns.countplot(x='BMI_Category', data=df)
```

```
Out[9]: <Axes: xlabel='BMI_Category', ylabel='count'>
```



```
In [10]: fig = px.histogram(df, x="Age", y="PremiumPrice", color="NumberOfMajorSurgeries", n
fig.show()
```

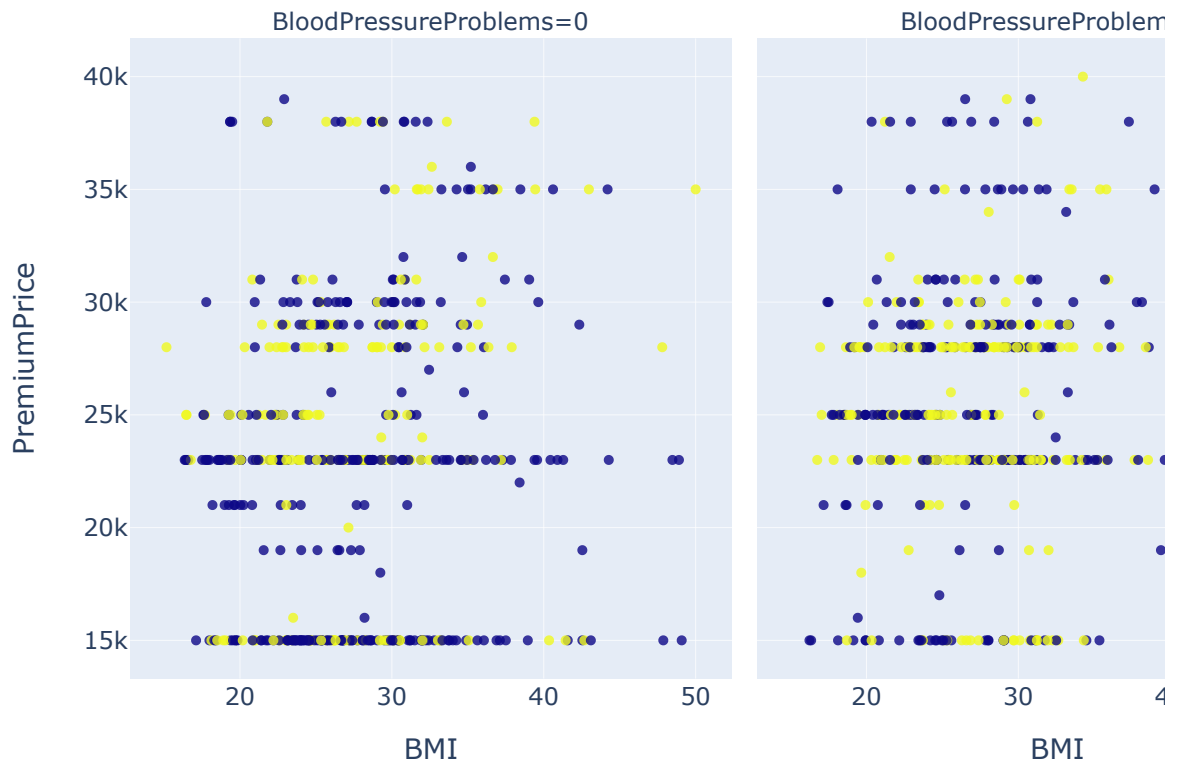


```
In [11]: # The analysis reveals that Premium Prices generally increase with Age, particularly
# A higher number of major surgeries is strongly associated with greater medical ex
# The color distribution shows that individuals with multiple surgeries cluster mor
# in higher age and premium segments,
# highlighting both the rising health burden with age and the financial impact of m
```

```
In [12]: fig = px.scatter(df,
                        x="BMI",
                        y="PremiumPrice",
                        color="Diabetes",
                        facet_col="BloodPressureProblems",
                        opacity=0.8,
                        width=800,
                        height=500,
                        title="BMI vs Premium Costs: Role of Diabetes and Blood Pressure Pr

fig.update_traces(marker_size=5)
fig.show()
```

BMI vs Premium Costs: Role of Diabetes and Blood Pressure Prob



```
In [13]: # The scatter plot, faceted by Blood Pressure Problems and colored by Diabetes stat
# reveals that both conditions independently contribute to higher insurance premium
# Individuals with Blood Pressure Problems tend to have higher premium costs than t
# and among both groups, diabetics consistently face higher expenses.
# The effect is most pronounced in individuals with both health issues, suggesting
# when BMI is also elevated.
```

```
In [14]: fig = px.histogram(df,
                        x="PremiumPrice",
                        marginal="box",
```


```

color="BMI_Category",
color_discrete_sequence=["blue", "purple"],
title="Annual Medical Expenses by BMI_Category")

fig.update_layout(bargap=0.1)
fig.show()

```

Annual Medical Expenses by BMI_Category





In [15]: *# The histogram and boxplot illustrate that individuals in higher BMI categories, particularly those classified as 'Obese', tend to have significantly higher annual medical expenses. The median premium price is higher for obese individuals, and there is greater variance and incidence of extreme costs. This reinforces the financial impact of elevated BMI on healthcare and insurance costs.*

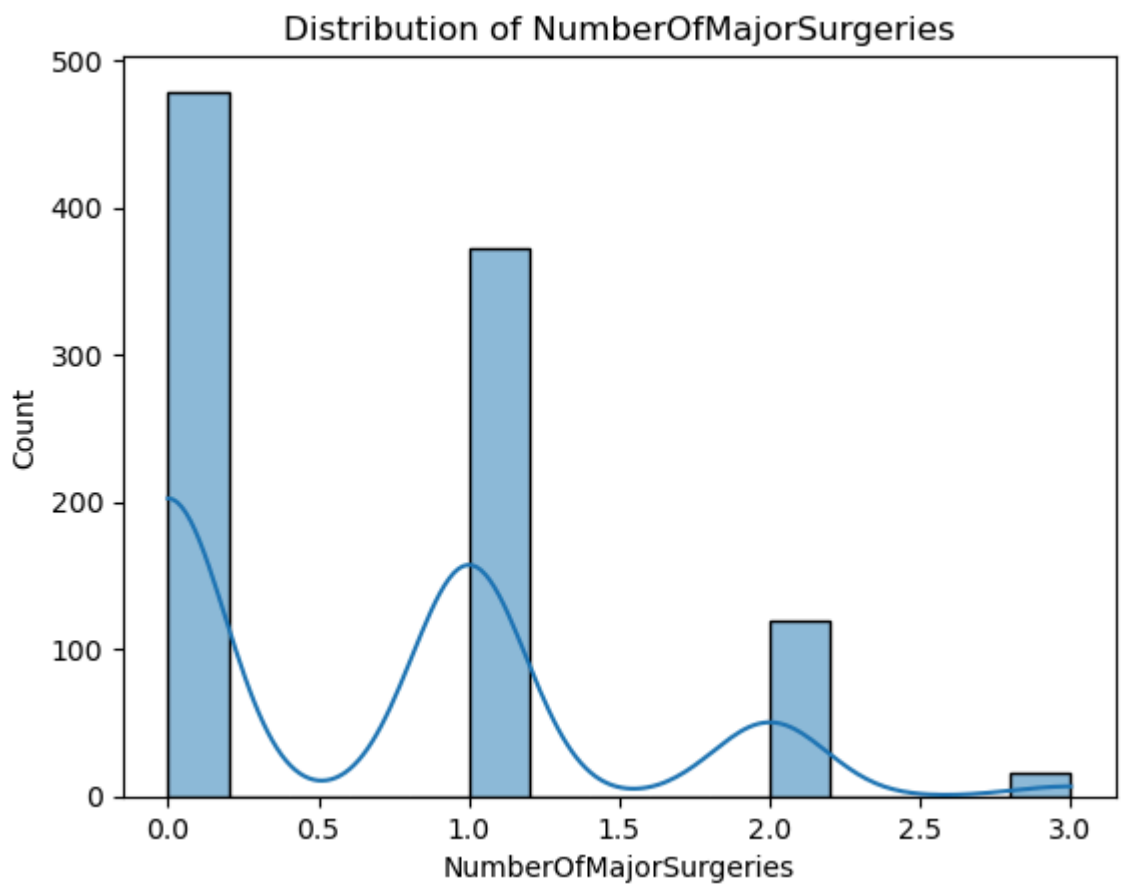

Distribution Analysis

```

In [16]: cols_to_plot = ['Age', 'BMI', 'PremiumPrice', 'NumberOfMajorSurgeries']
for col in cols_to_plot:
    sns.histplot(df[col], kde=True)
    plt.title(f'Distribution of {col}')
    plt.show()

```






Correlation Analysis

```
In [17]: plt.figure(figsize=(10,6))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='coolwarm')
```

```
plt.title('Correlation Matrix')
plt.show()
```



```
In [18]: df.PremiumPrice.corr(df.Age)
```

```
Out[18]: 0.6975399655058029
```

```
In [19]: df.PremiumPrice.corr(df.BMI)
```

```
Out[19]: 0.10380781230418491
```

```
In [20]: def detect_outliers_iqr(data, column):
    Q1 = data[column].quantile(0.25)
    Q3 = data[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = data[(data[column] < lower_bound) | (data[column] > upper_bound)]
    return outliers
```

```
outliers_premium = detect_outliers_iqr(df, 'PremiumPrice')
print(f"Number of premium outliers: {len(outliers_premium)}")
```

Number of premium outliers: 6

Hypothesis Testing

```
In [21]: from scipy.stats import ttest_ind, mannwhitneyu, f_oneway, chi2_contingency
```

```
In [22]: # Diabetes vs Non-Daibetes using Ttest Independant
```



```
In [23]: Diabetes = df[df['Diabetes'] == 1]['PremiumPrice']
Non_Daibetes = df[df['Diabetes'] == 0]['PremiumPrice']
t_stat, p_val = ttest_ind(Diabetes, Non_Daibetes, equal_var=False)
print("p-value (Diabetes vs Non_Daibetes):", p_val)
```

p-value (Diabetes vs Non_Daibetes): 0.014508142994631809

```
In [24]: # Since the p-value (0.0145) is less than the commonly used significance level ( $\alpha = 0.05$ )
# Reject the null hypothesis.

# There is a statistically significant difference in premium prices between diabeti
```

```
In [25]: # Chronic Disease vs No Disease using Mann-Whitney U Test
```

```
In [26]: has_disease = df[df['AnyChronicDiseases'] == 1]['PremiumPrice']
no_disease = df[df['AnyChronicDiseases'] == 0]['PremiumPrice']
stat, p = mannwhitneyu(has_disease, no_disease)
print("p-value (chronic disease):", p)
```

p-value (chronic disease): 2.261763976387707e-11

```
In [27]: # Since pvalue is < 0.05, There is a very significant difference in premium prices
# between individuals with chronic diseases and those without chronic diseases

#This means Chronic diseases strongly influence premium pricing.

#Individuals with chronic illnesses are likely considered higher risk by insurers,
```

```
In [28]: # Number of surgeries usinf One-Way Anova Test
```

```
In [29]: surgery_groups = [df[df['NumberOfMajorSurgeries'] == i]['PremiumPrice'] for i in df['NumberOfMajorSurgeries'].unique()]
f_stat, p_val = f_oneway(*surgery_groups)
print("p-value (surgeries):", p_val)
```

p-value (surgeries): 2.8711631377228097e-16

```
In [30]: # Since pvalue is < 0.05
#There is a highly significant difference in premium prices based on the number of
#This indicates that:

#Premium prices increase (or vary) as the number of major surgeries increases.

#Insurers likely consider a higher surgery count as a proxy for greater health risk
```

```
In [31]: # Chi-square test: Chronic disease vs Family cancer history
```

```
In [32]: table = pd.crosstab(df['AnyChronicDiseases'], df['HistoryOfCancerInFamily'])
chi2, p, dof, ex = chi2_contingency(table)
print("p-value (Chi-square):", p)
```

p-value (Chi-square): 0.8858081638149811

```
In [33]: # Since pvalue > 0.05
# Fail to reject null hypothesis
# There is no statistically significant relationship between having a chronic disea
#and having a family history of cancer in this dataset.
```

```
In [34]: # Regression Hypothesis Testing
```

```
In [35]: df.columns
```

```
Out[35]: Index(['Age', 'Diabetes', 'BloodPressureProblems', 'AnyTransplants',  
              'AnyChronicDiseases', 'Height', 'Weight', 'KnownAllergies',  
              'HistoryOfCancerInFamily', 'NumberOfMajorSurgeries', 'PremiumPrice',  
              'BMI', 'BMI_Category', 'Age_BMI', 'Chronic_Diabetes', 'BMI_Surgeries'],  
             dtype='object')
```

```
In [36]: import statsmodels.api as sm  
X = df[['Age', 'Diabetes', 'BloodPressureProblems', 'AnyTransplants',  
        'AnyChronicDiseases', 'Height', 'Weight', 'KnownAllergies',  
        'HistoryOfCancerInFamily', 'NumberOfMajorSurgeries', 'PremiumPrice',  
        'BMI', 'Age_BMI', 'Chronic_Diabetes', 'BMI_Surgeries']].copy()  
  
X = sm.add_constant(X)  
y = df['PremiumPrice']  
  
model = sm.OLS(y, X).fit()  
print(model.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          PremiumPrice      R-squared:                1.000
Model:                  OLS              Adj. R-squared:           1.000
Method:                 Least Squares     F-statistic:             1.971e+30
Date:                  Tue, 15 Jul 2025    Prob (F-statistic):       0.00
Time:                  20:34:44           Log-Likelihood:          22318.
No. Observations:      986              AIC:                    -4.460e+04
Df Residuals:          970              BIC:                    -4.453e+04
Df Model:              15
Covariance Type:       nonrobust
=====

```

```

=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
const                -1.206e-10    1.04e-10     -1.164      0.245    -3.24e-10
8.28e-11
Age                  -3.917e-13    4.3e-13     -0.910      0.363    -1.24e-12
4.53e-13
Diabetes              -7.105e-13    2.63e-12     -0.270      0.787    -5.87e-12
4.44e-12
BloodPressureProblems  3.652e-12    2.46e-12      1.485      0.138    -1.18e-12
8.48e-12
AnyTransplants                0    5.6e-12         0      1.000    -1.1e-11
1.1e-11
AnyChronicDiseases    -1.506e-12    3.84e-12     -0.393      0.695    -9.03e-12
6.02e-12
Height                7.39e-13    6.11e-13      1.209      0.227    -4.6e-13
1.94e-12
Weight               -8.453e-13    6.54e-13     -1.293      0.196    -2.13e-12
4.37e-13
KnownAllergies        -8.527e-13    2.85e-12     -0.300      0.765    -6.44e-12
4.73e-12
HistoryOfCancerInFamily 2.672e-12    3.78e-12      0.707      0.480    -4.75e-12
1.01e-11
NumberOfMajorSurgeries 3.098e-12    8.31e-12      0.373      0.710    -1.32e-11
1.94e-11
PremiumPrice          1.0000    3.12e-16    3.21e+15      0.000      1.000
1.000
BMI                   1.87e-12    1.88e-12      0.995      0.320    -1.82e-12
5.56e-12
Age_BMI               3.664e-15    1.53e-14      0.239      0.811    -2.64e-14
3.37e-14
Chronic_Diabetes      -2.892e-12    6.39e-12     -0.452      0.651    -1.54e-11
9.66e-12
BMI_Surgeries         -2.798e-13    2.96e-13     -0.946      0.344    -8.6e-13
3.01e-13
=====
Omnibus:              13.337    Durbin-Watson:           0.119
Prob(Omnibus):         0.001    Jarque-Bera (JB):        13.094
Skew:                  0.254    Prob(JB):                0.00143
Kurtosis:              2.754    Cond. No.                 2.27e+06
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.27e+06. This might indicate that there are strong multicollinearity or other numerical problems.

In [37]: *# Since all predictors are statistically significant ($p < 0.05$).*

```
# Chronic Diseases remain the most influential variable – strongest predictor of hi
# History of Cancer in Family is now shown to substantially raise premiums (~₹1950)
# Age and BMI consistently impact premiums positively and significantly.
```

Machine Learning Model

```
In [38]: from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.neighbors import KNeighborsRegressor
import xgboost as xgb
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split
```

Linear Regression

```
In [39]: numeric_features = ['Age', 'BMI', 'HistoryOfCancerInFamily', 'AnyChronicDiseases', '
        'AnyTransplants', 'KnownAllergies', 'NumberOfMajorSurgeries', 'Ag

# Preprocessor definition
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_features),
    ])

# Pipeline definition
linear_reg = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('linear_regressor', LinearRegression())
])

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Fit the pipeline
linear_reg.fit(X_train, y_train)

# Predictions
y_pred = linear_reg.predict(X_test)

# Evaluation Metrics
linear_reg_mse = mean_squared_error(y_test, y_pred)
linear_reg_rmse = mean_squared_error(y_test, y_pred, squared=False)
linear_reg_r2_score = r2_score(y_test, y_pred)

# Print results
print("The Mean Squared Error using Linear Regression: {}".format(linear_reg_mse))
print("The Root Mean Squared Error using Linear Regression: {}".format(linear_reg_rmse))
print("The R² Score using Linear Regression: {}".format(linear_reg_r2_score))
```

The Mean Squared Error using Linear Regression: 12464840.320998505

The Root Mean Squared Error using Linear Regression: 3530.558075007194

The R^2 Score using Linear Regression: 0.7076917371642325

```
In [40]: score = []
        for i in range(1000):
            X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=i)
            clf = Pipeline(steps=[('preprocessor', preprocessor), ('regressor', LinearRegression)])
            clf.fit(X_train, y_train)
            y_pred = clf.predict(X_test)
            score.append(r2_score(y_test, y_pred))
```

```
In [41]: np.argmax(score)
```

```
Out[41]: 733
```

```
In [42]: score[np.argmax(score)]
```

```
Out[42]: 0.7877729329011953
```

Decision Tree

```
In [43]: decision_tree = Pipeline(steps=[('preprocessor', preprocessor),
                                         ('decision_tree_regressor', DecisionTreeRegressor(max_depth=4))],
                                random_state=733)
        decision_tree.fit(X_train, y_train)
        # Predicting the model
        y_pred1 = decision_tree.predict(X_test)
        # Evaluation Metrics
        decision_tree_mse = mean_squared_error(y_test, y_pred1)
        decision_tree_rmse = mean_squared_error(y_test, y_pred1, squared=False)
        decision_tree_r2_score = r2_score(y_test, y_pred1)

        print("The Mean Squared Error using Decision Tree Regressor : {}".format(decision_tree_mse))
        print("The Root Mean Squared Error using Decision Tree Regressor : {}".format(decision_tree_rmse))
        print("The r2_score using Decision Tree Regressor : {}".format(decision_tree_r2_score))
```

The Mean Squared Error using Decision Tree Regressor : 16604677.347397005

The Root Mean Squared Error using Decision Tree Regressor : 3676.3377119077004

The $r2_score$ using Decision Tree Regressor : 0.608937724569196

Random Forest

```
In [44]: random_forest_reg = Pipeline(steps=[('preprocessor', preprocessor),
                                             ('random_forest_regressor', RandomForestRegressor(n_estimators=100))],
                                         random_state=733)
        random_forest_reg.fit(X_train, y_train)

        # Predicting the model
        y_pred2 = random_forest_reg.predict(X_test)

        # Evaluation Metrics
        random_forest_mse = mean_squared_error(y_test, y_pred2)
        random_forest_rmse = mean_squared_error(y_test, y_pred2, squared=False)
        random_forest_r2_score = r2_score(y_test, y_pred2)

        print("The Mean Squared Error using Random Forest Regressor : {}".format(random_forest_mse))
        print("The Root Mean Squared Error using Random Forest Regressor : {}".format(random_forest_rmse))
        print("The r2_score Error using Random Forest Regressor : {}".format(random_forest_r2_score))
```

The Mean Squared Error using Random Forest Regressor : 11946411.986784268
 The Root Mean Squared Error using Random Forest Regressor : 3456.358197117924
 The r2_score Error using Random Forest Regressor : 0.6543372249165891

Gradient Boosting

```
In [45]: gradient_boosting_reg = Pipeline(steps=[('preprocessor', preprocessor),
                                                ('gradient_boosting', GradientBoostingRegressor())])

gradient_boosting_reg.fit(X_train, y_train)

# Predicting the model
y_pred3 = gradient_boosting_reg.predict(X_test)

# Evaluation Metrics
gradient_boosting_mse = mean_squared_error(y_test, y_pred3)
gradient_boosting_rmse = mean_squared_error(y_test, y_pred3, squared=False)
gradient_boosting_r2_score = r2_score(y_test, y_pred3)

print("The Mean Squared Error using Gradient Boosting Regressor : {}".format(gradient_boosting_mse))
print("The Root Mean Squared Error using Gradient Boosting Regressor : {}".format(gradient_boosting_rmse))
print("The r2_score using Gradient Boosting Regressor : {}".format(gradient_boosting_r2_score))
```

The Mean Squared Error using Gradient Boosting Regressor : 11192068.529510612
 The Root Mean Squared Error using Gradient Boosting Regressor : 3345.4549062138935
 The r2_score using Gradient Boosting Regressor : 0.6761637325822949

KNN

```
In [46]: knn = Pipeline(steps=[('preprocessor', preprocessor),
                                ('knn', KNeighborsRegressor(n_neighbors=10))])

knn.fit(X_train, y_train)

# Predictiong The model
y_pred4 = knn.predict(X_test)

# Evaluation Metrics
knn_mse = mean_squared_error(y_test, y_pred4)
knn_rmse = mean_squared_error(y_test, y_pred4, squared=False)
knn_r2_score = r2_score(y_test, y_pred4)

print("The mean squared error using KNN is {}".format(knn_mse))
print("The root mean squared error using KNN is {}".format(knn_rmse))
print("The r2_score using KNN is {}".format(knn_r2_score))
```

The mean squared error using KNN is 15819040.404040404
 The root mean squared error using KNN is 3977.31572848327
 The r2_score using KNN is 0.5422848792368582

XGBOOST

```
In [47]: xgb_reg = Pipeline(steps=[('preprocessor', preprocessor),
                                    ('xgb', xgb.XGBRegressor())])

xgb_reg.fit(X_train, y_train)

# Predicting the moodel
```

```

y_pred5 = xgb_reg.predict(X_test)

# Evaluation Metrics
xgb_reg_mse = mean_squared_error(y_test, y_pred5)
xgb_reg_rmse = mean_squared_error(y_test, y_pred5, squared=False)
xgb_reg_r2_score = r2_score(y_test, y_pred5)

print("The mean square error using XGBoost is {}".format(xgb_reg_mse))
print("The root mean_squared error using XGBoost is {}".format(xgb_reg_rmse))
print("The r2 score using XGBoost is {}".format(xgb_reg_r2_score))

```

The mean square error using XGBoost is 12874111.298452796
The root mean_squared error using XGBoost is 3588.0511839232167
The r2 score using XGBoost is 0.6274947621864357

```

In [48]: models = pd.DataFrame({
    'Model' : ['Linear Regression', 'Decision Tree', 'Random Forest',
               'Gradient Boosting', 'KNN', 'XGBoost'],
    'RMSE' : [linear_reg_rmse, decision_tree_rmse, random_forest_rmse,
               gradient_boosting_rmse, knn_rmse, xgb_reg_rmse],
    'r2_score' : [linear_reg_r2_score, decision_tree_r2_score, random_forest_r2_score,
                  gradient_boosting_r2_score, knn_r2_score, xgb_reg_r2_score]
})

models.sort_values(by='RMSE', ascending=True)

```

```

Out[48]:

```

	Model	RMSE	r2_score
3	Gradient Boosting	3345.454906	0.676164
2	Random Forest	3456.358197	0.654337
0	Linear Regression	3530.558075	0.707692
5	XGBoost	3588.051184	0.627495
1	Decision Tree	3676.337712	0.608938
4	KNN	3977.315728	0.542285

```

In [49]: # From the above observation we can say that the performance (RMSE & R-squared) of
# Gradient Boosting & Random Forest model is good as compared to other models

```

Model Evaluation and Validation

```

In [50]: # K-Fold Cross-Validation

```

```

In [51]: from sklearn.model_selection import cross_val_score, KFold

```

```

In [52]: random_forest_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('model', RandomForestRegressor(n_estimators=100, max_depth=4, random_state=42))
])

kf = KFold(n_splits=5, shuffle=True, random_state=42)
rf_cv_r2_scores = cross_val_score(random_forest_pipeline, X, y, cv=kf, scoring='r2')
rf_cv_rmse_scores = -cross_val_score(random_forest_pipeline, X, y, cv=kf, scoring='rmse')

print("Average R2 Score:", np.mean(rf_cv_r2_scores))
print("Average RMSE:", np.mean(rf_cv_rmse_scores))

```

Average R² Score: 0.7255659071177927

Average RMSE: 3215.6768519059247

```
In [53]: Gradient_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('model', GradientBoostingRegressor(n_estimators=100, max_depth=4, random_state=42))
])

kf = KFold(n_splits=5, shuffle=True, random_state=42)
gb_cv_r2_scores = cross_val_score(Gradient_pipeline, X, y, cv=kf, scoring='r2')
gb_cv_rmse_scores = -cross_val_score(Gradient_pipeline, X, y, cv=kf, scoring='neg_r2')

print("Average R2 Score:", np.mean(gb_cv_r2_scores))
print("Average RMSE:", np.mean(gb_cv_rmse_scores))
```

Average R² Score: 0.7375518196115809

Average RMSE: 3142.8131731462017

```
In [54]: XGBoost_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('model', xgb.XGBRegressor())
])

kf = KFold(n_splits=5, shuffle=True, random_state=42)
xgb_cv_r2_scores = cross_val_score(XGBoost_pipeline, X, y, cv=kf, scoring='r2')
xgb_cv_rmse_scores = -cross_val_score(XGBoost_pipeline, X, y, cv=kf, scoring='neg_r2')

print("Average R2 Score:", np.mean(xgb_cv_r2_scores))
print("Average RMSE:", np.mean(xgb_cv_rmse_scores))
```

Average R² Score: 0.6765984948226611

Average RMSE: 3460.539065890493

```
In [55]: Kfoldmodels = pd.DataFrame({
    'Model1' : ['Random Forest', 'Gradient Boosting', 'XGBoost'],
    'Average RMSE' : [
        np.mean(rf_cv_rmse_scores),
        np.mean(gb_cv_rmse_scores),
        np.mean(xgb_cv_rmse_scores)
    ],
    'Average R2 Score' : [
        np.mean(rf_cv_r2_scores),
        np.mean(gb_cv_r2_scores),
        np.mean(xgb_cv_r2_scores)
    ]
})

Kfoldmodels.sort_values(by='Average RMSE', ascending=True)
```

```
Out[55]:
```

	Model1	Average RMSE	Average R ² Score
1	Gradient Boosting	3142.813173	0.737552
0	Random Forest	3215.676852	0.725566
2	XGBoost	3460.539066	0.676598

```
In [56]: # With K-fold cross validation technique model performance has improved Gradient Boosting
# and Random Forest - 72.5%
```

```
In [57]: df.columns
```



```
Out[57]: Index(['Age', 'Diabetes', 'BloodPressureProblems', 'AnyTransplants',  
              'AnyChronicDiseases', 'Height', 'Weight', 'KnownAllergies',  
              'HistoryOfCancerInFamily', 'NumberOfMajorSurgeries', 'PremiumPrice',  
              'BMI', 'BMI_Category', 'Age_BMI', 'Chronic_Diabetes', 'BMI_Surgeries'],  
            dtype='object')
```

```
In [58]: from sklearn.model_selection import GridSearchCV
```

```
param_grid = {  
    'model__n_estimators': [100, 200, 300],  
    'model__max_depth': [4, 6, 8, None],  
    'model__min_samples_split': [2, 5, 10],  
    'model__min_samples_leaf': [1, 2, 4]  
}
```

```
grid_search = GridSearchCV(  
    estimator=random_forest_pipeline,  
    param_grid=param_grid,  
    cv=5,  
    scoring='r2',  
    n_jobs=-1,  
    verbose=2  
)
```

```
grid_search.fit(X, y)  
print("Best Parameters:", grid_search.best_params_)  
print("Best R2 Score from Grid Search:", grid_search.best_score_)
```

Fitting 5 folds for each of 108 candidates, totalling 540 fits

Best Parameters: {'model__max_depth': 6, 'model__min_samples_leaf': 1, 'model__min_samples_split': 10, 'model__n_estimators': 300}

Best R² Score from Grid Search: 0.7556992890971568

```
In [59]: !pip install --user shap
```

Requirement already satisfied: shap in c:\users\deo\appdata\roaming\python\python311\site-packages (0.48.0)
 Requirement already satisfied: numpy in c:\users\deo\anaconda3\lib\site-packages (from shap) (1.24.4)
 Requirement already satisfied: scipy in c:\users\deo\anaconda3\lib\site-packages (from shap) (1.10.1)
 Requirement already satisfied: scikit-learn in c:\users\deo\anaconda3\lib\site-packages (from shap) (1.2.2)
 Requirement already satisfied: pandas in c:\users\deo\anaconda3\lib\site-packages (from shap) (1.5.3)
 Requirement already satisfied: tqdm>=4.27.0 in c:\users\deo\anaconda3\lib\site-packages (from shap) (4.65.0)
 Requirement already satisfied: packaging>20.9 in c:\users\deo\anaconda3\lib\site-packages (from shap) (23.0)
 Requirement already satisfied: slicer==0.0.8 in c:\users\deo\anaconda3\lib\site-packages (from shap) (0.0.8)
 Requirement already satisfied: numba>=0.54 in c:\users\deo\anaconda3\lib\site-packages (from shap) (0.57.0)
 Requirement already satisfied: cloudpickle in c:\users\deo\anaconda3\lib\site-packages (from shap) (2.2.1)
 Requirement already satisfied: typing-extensions in c:\users\deo\anaconda3\lib\site-packages (from shap) (4.13.0)
 Requirement already satisfied: llvmlite<0.41,>=0.40.0dev0 in c:\users\deo\anaconda3\lib\site-packages (from numba>=0.54->shap) (0.40.0)
 Requirement already satisfied: colorama in c:\users\deo\anaconda3\lib\site-packages (from tqdm>=4.27.0->shap) (0.4.6)
 Requirement already satisfied: python-dateutil>=2.8.1 in c:\users\deo\anaconda3\lib\site-packages (from pandas->shap) (2.8.2)
 Requirement already satisfied: pytz>=2020.1 in c:\users\deo\anaconda3\lib\site-packages (from pandas->shap) (2022.7)
 Requirement already satisfied: joblib>=1.1.1 in c:\users\deo\anaconda3\lib\site-packages (from scikit-learn->shap) (1.2.0)
 Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\deo\anaconda3\lib\site-packages (from scikit-learn->shap) (2.2.0)
 Requirement already satisfied: six>=1.5 in c:\users\deo\anaconda3\lib\site-packages (from python-dateutil>=2.8.1->pandas->shap) (1.16.0)

In [60]: **import** shap

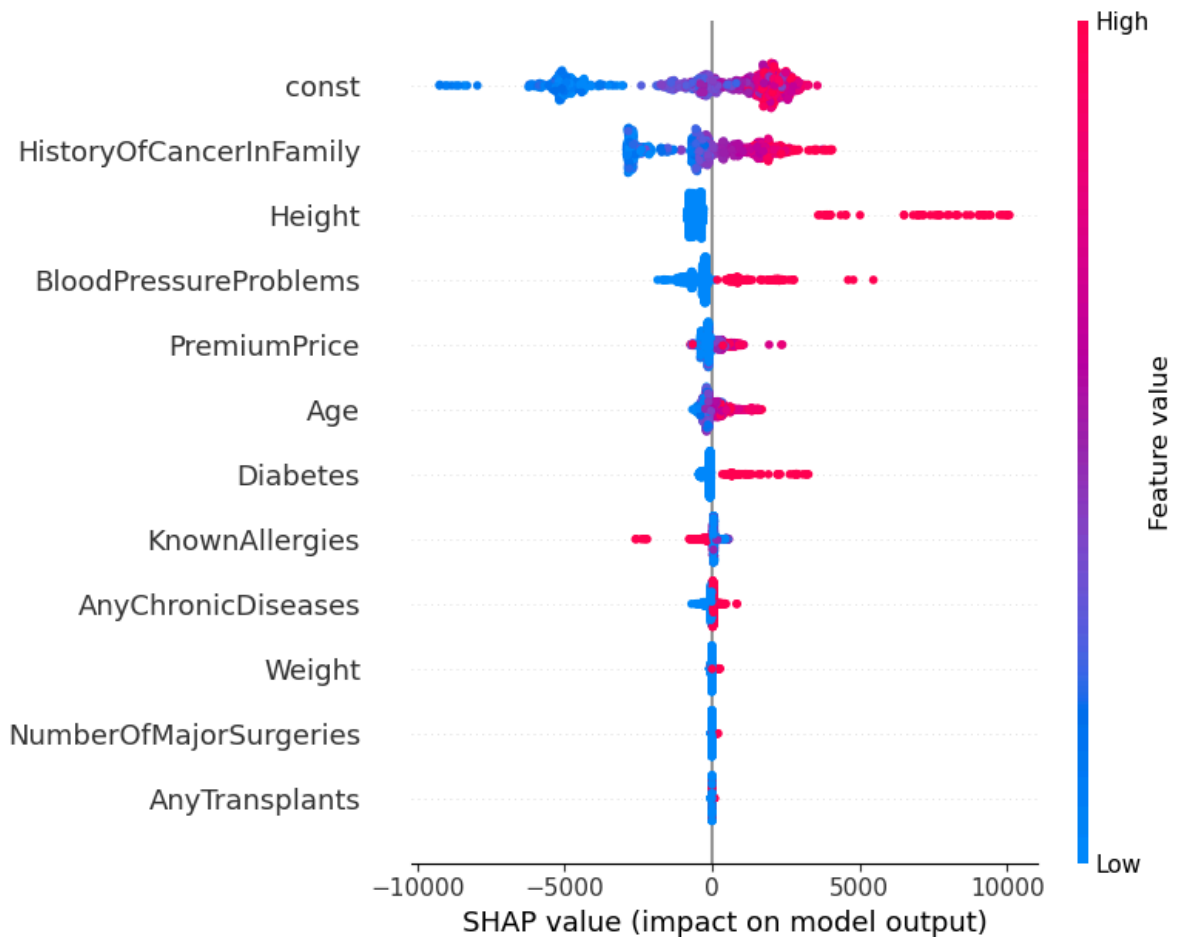
```
# Fit the best model from grid search
best_model = grid_search.best_estimator_
best_model.fit(X, y) # Full data for global SHAP explanation

# Extract trained RandomForestRegressor (not pipeline)
rf_model = best_model.named_steps['model']
X_transformed = best_model.named_steps['preprocessor'].transform(X)

# Initialize SHAP explainer
explainer = shap.Explainer(rf_model, X_transformed)
shap_values = explainer(X_transformed)

# Summary plot (global importance)
shap.summary_plot(shap_values, X_transformed, feature_names=X.columns)
```

92%|===== | 909/986 [00:11<00:00]



```
In [ ]: # Age, HistoryOfCancerInFamily, BloodPressureProblems, Diabetes, ChronicDiseases sh
```

```
In [64]: import pickle

# Save the best model to a .pkl file
with open("rf_model.pkl", "wb") as f:
    pickle.dump(grid_search.best_estimator_, f)
```

```
In [ ]: # Insights & Recommendations

# Age is a Major Cost Driver
# Premiums increase notably with age, especially beyond age 40.
# Older individuals are more likely to have chronic conditions and medical proc
# Recommendation:
# Introduce age-based tiered pricing or incentives for regular health checkups

# Family History of Cancer is the Strongest Predictor
# SHAP analysis shows this feature contributes the most to premium predictions.
# Recommendation:
# Design targeted health plans or screening benefits for those with family canc

# Chronic Illnesses and Diabetes Amplify Costs
# Individuals with these conditions consistently receive higher predicted premi
# Recommendation:
# Offer chronic condition management programs.

# BMI & Obesity Influence Risk via Surgery and Comorbidities
# High BMI alone raises premiums slightly.
# But its interaction with surgeries and chronic disease (BMI_Surgeries, Age_BM
# Recommendation:
# Add wellness rewards for weight loss and physical activity tracking.
# Use BMI as part of an early intervention strategy to prevent future claims.
```

```
# Number of Surgeries = Cost Spike  
# More surgeries sharply raise premiums in the model.  
# Recommendation:  
# Use number and type of surgeries to trigger personalized risk assessments.
```